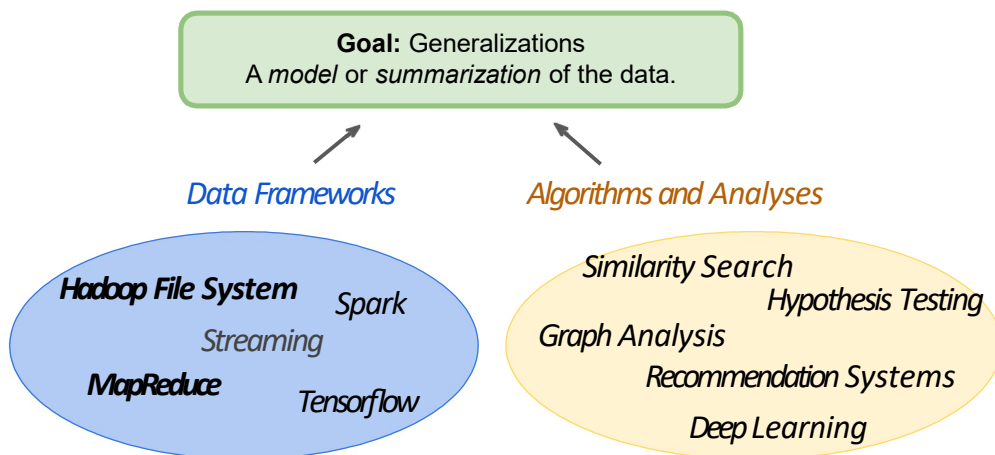


“Hadoop”:

A Distributed Architecture, FileSystem, & MapReduce


1

Big Data Analytics, The Class




2

Big Data Analytics, The Class




3

Big Data Analytics, The Class




4

Classical Data Analytics




5

Classical Data Analytics




6

Classical Data Analytics



7

Classical Data Analytics




8

IO Bounded

Reading a word from disk versus main memory: 10^5 slower!

Reading many contiguously stored words is faster per word, but fast modern disks still only reach 150MB/s for sequential reads.




9

IO Bounded

Reading a word from disk versus main memory: 10^5 slower!

Reading many contiguously stored words is faster per word, but fast modern disks still only reach 150MB/s for sequential reads.




IO Bound: biggest performance bottleneck is reading / writing to disk.

starts around 100 GBs: ~10 minutes just to read

200 TBs: ~20,000 minutes = 13 days

10

Classical Big Data



Classical focus: efficient use of disk.
e.g. Apache Lucene / Solr

Classical limitation: Still bounded when needing to process all of a large file.

11


Classical Big Data

How to solve?

Classical limitation: Still bounded when needing to process all of a large file.

12

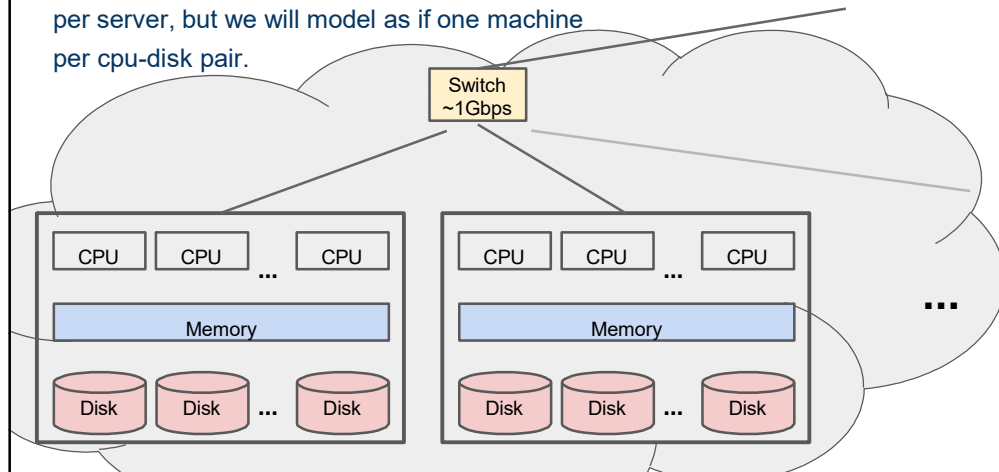
Distributed Architecture



13


Distributed Architecture

In reality, modern setups often have multiple cpus and disks per server, but we will model as if one machine per cpu-disk pair.



14

Distributed Architecture (Cluster)



15

Distributed Architecture (Cluster)

Challenges for IO Cluster Computing

1. Nodes fail
1 in 1000 nodes fail a day
2. Network is a bottleneck
Typically 1-10 Gb/s throughput
3. Traditional distributed programming is
often ad-hoc and complicated

16

Distributed Architecture (Cluster)

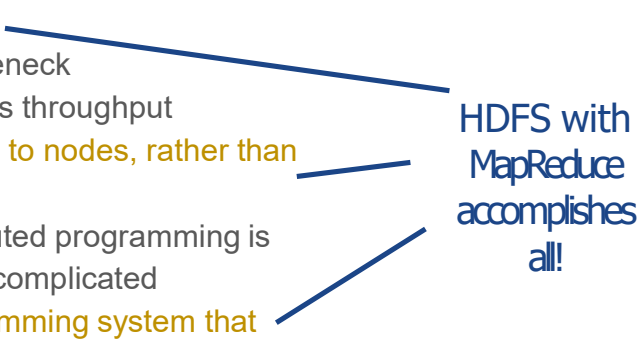
Challenges for IO Cluster Computing

1. Nodes fail
1 in 1000 nodes fail a day
Duplicate Data
2. Network is a bottleneck
Typically 1-10 Gb/s throughput
Bring computation to nodes, rather than data to nodes.
3. Traditional distributed programming is often ad-hoc and complicated
Stipulate a programming system that can easily be distributed

17

Distributed Architecture (Cluster)

Challenges for IO Cluster Computing

1. Nodes fail
1 in 1000 nodes fail a day
Duplicate Data
 2. Network is a bottleneck
Typically 1-10 Gb/s throughput
Bring computation to nodes, rather than data to nodes.
 3. Traditional distributed programming is often ad-hoc and complicated
Stipulate a programming system that can easily be distributed
- 
- HDFS with MapReduce accomplishes all!

18

Distributed Filesystem

The effectiveness of MapReduce, Spark, and other distributed processing systems is in part simply due to use of a distributed filesystem!

19



Distributed Filesystem

Characteristics for Big Data Tasks

Large files (i.e. >100 GB to TBs)

Reads are most common

No need to update in place
(append preferred)



20

Distributed Filesystem

(e.g. Apache HadoopDFS, GoogleFS, EMRFS)

C, D: Two different files




<https://opensource.com/life/14/8/intro-apache-hadoop-big-data>

21

Distributed Filesystem

(e.g. Apache **Hadoop**DFS, GoogleFS, EMRFS)

C, D: Two different files



<https://opensource.com/life/14/8/intro-apache-hadoop-big-data>


"Hadoop" was named after a toy elephant belonging to Doug Cutting's son. Cutting was one of Hadoop's creators.

22

Distributed Filesystem

(e.g. Apache HadoopDFS, GoogleFS, EMRFS)

C, D: Two different files; break into chunks (or "partitions"):




23

Distributed Filesystem

(e.g. Apache HadoopDFS, GoogleFS, EMRFS)

C, D: Two different files



(Leskovec et al., 2014; <http://www.mmms.org/>)

24

Distributed Filesystem

(e.g. Apache HadoopDFS, GoogleFS, EMRFS)

C, D: Two different files




(Leskovec et al., 2014; <http://www.mmds.org/>)

25

Distributed Filesystem

(e.g. Apache HadoopDFS, GoogleFS, EMRFS)

C, D: Two different files



(Leskovec et al., 2014; <http://www.mmds.org/>)

26

Distributed Filesystem

Chunk servers (on Data Nodes)

- File is split into contiguous chunks
- Typically each chunk is 16-64MB
- Each chunk replicated (usually 2x or 3x)
- Try to keep replicas in different racks

(Leskovec et al., 2014; <http://www.mmds.org/>)

27

Components of a Distributed Filesystem

Chunk servers (on Data Nodes)

- File is split into contiguous chunks
- Typically each chunk is 16-64MB
- Each chunk replicated (usually 2x or 3x)
- Try to keep replicas in different racks

Name node (aka master node)

- Stores metadata about where files are stored
- Might be replicated or distributed across data nodes.

(Leskovec et al., 2014; <http://www.mmds.org/>)

28

Components of a Distributed Filesystem

Chunk servers (on Data Nodes)

- File is split into contiguous chunks
- Typically each chunk is 16-64MB
- Each chunk replicated (usually 2x or 3x)
- Try to keep replicas in different racks

Name node (aka master node)

- Stores metadata about where files are stored
- Might be replicated or distributed across data nodes.

Client library for file access


- Talks to master to find chunk servers
- Connects directly to chunk servers to access data

(Leskovec et al., 2014; <http://www.mmds.org/>)

29

Distributed Architecture (Cluster)

Challenges for IO Cluster Computing

1. Nodes fail
1 in 1000 nodes fail a day
Duplicate Data (Distributed FS) 
2. Network is a bottleneck
Typically 1-10 Gb/s throughput
Bring computation to nodes, rather than data to nodes.
3. Traditional distributed programming is often ad-hoc and complicated
Stipulate a programming system that can easily be distributed

30

What is MapReduce

noun.1 - A style of programming

input chunks => **map tasks** | group_by keys | **reduce tasks** => output

“|” is the linux “pipe” symbol: passes stdout from first process to stdin of next.

31

What is MapReduce

noun.1 - A style of programming

input chunks => **map tasks** | group_by keys | **reduce tasks** => output

“|” is the linux “pipe” symbol: passes stdout from first process to stdin of next.

E.g. counting words:

```
tokenize(document) | sort | uniq -c
```

32

What is MapReduce

noun.1 - A style of programming

input chunks => **map tasks** | group_by keys | **reduce tasks** => output

"|" is the linux "pipe" symbol: passes stdout from first process to stdin of next.

E.g. counting words:

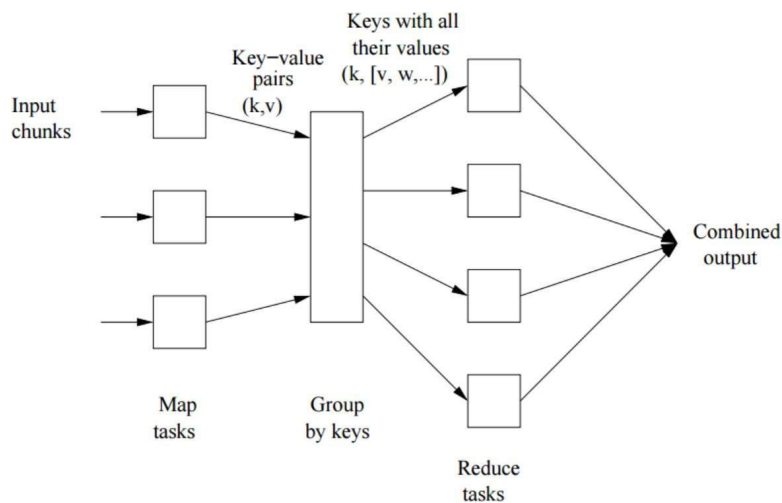
```
tokenize(document) | sort | uniq -c
```

noun.2 - A system that distributes MapReduce style programs across a distributed file-system.

(e.g. Google's internal "MapReduce" or apache.hadoop.mapreduce with hdfs)

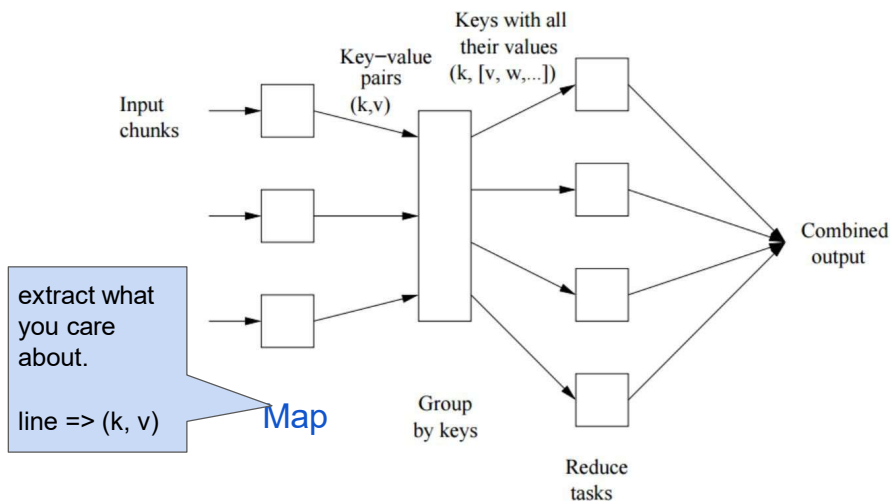
33

What is MapReduce



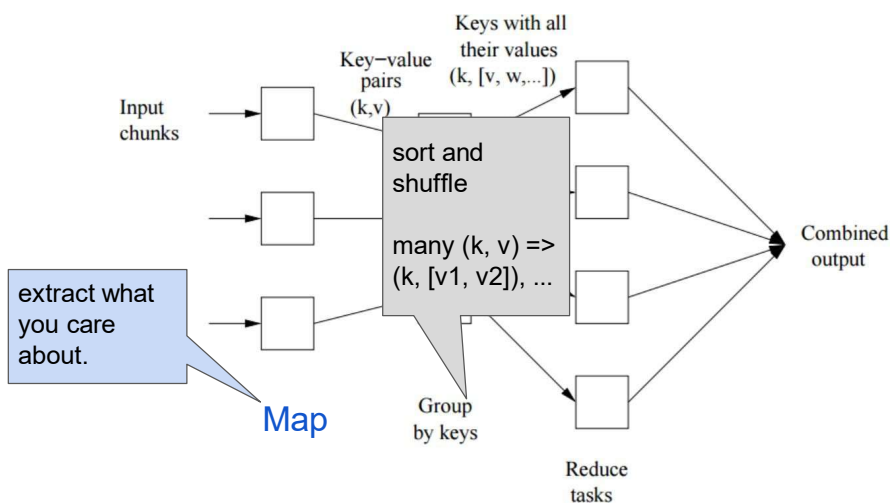
34

What is MapReduce



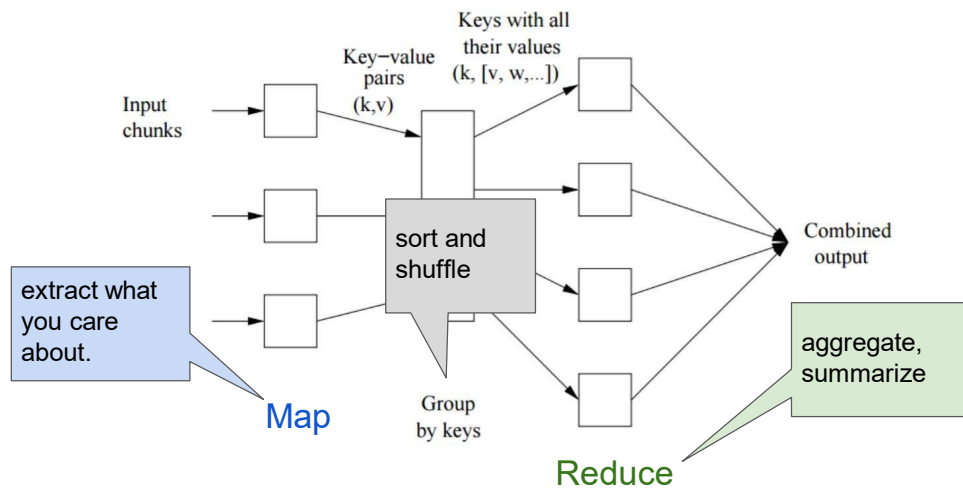
35

What is MapReduce



36

What is MapReduce



37

What is MapReduce

Easy as 1, 2, 3!

*Step 1: **Map***

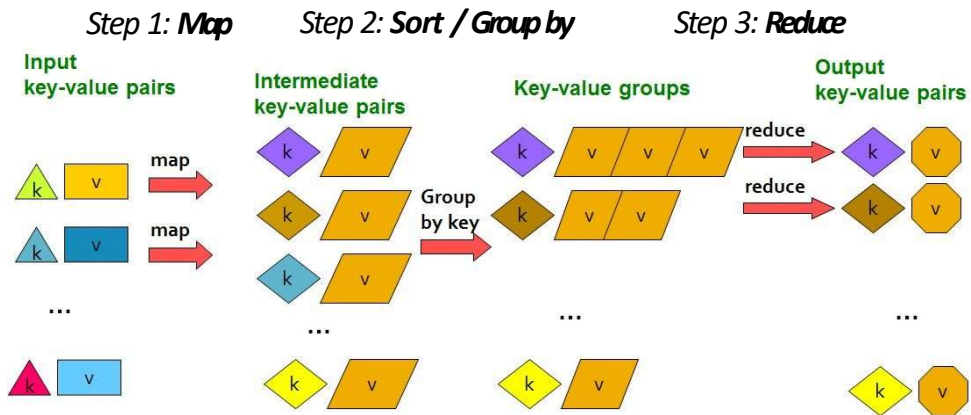
*Step 2: **Sort / Group by***

*Step 3: **Reduce***

38

What is MapReduce


Easy as 1, 2, 3!



(Leskovec et al., 2014; <http://www.mmds.org/>)

39


(1) The Map Step



(Leskovec et al., 2014; <http://www.mmds.org/>)

40


(2) The Sort / Group-by Step



(Leskovec et al., 2014; <http://www.mmds.org/>)

41

(3) The Reduce Step

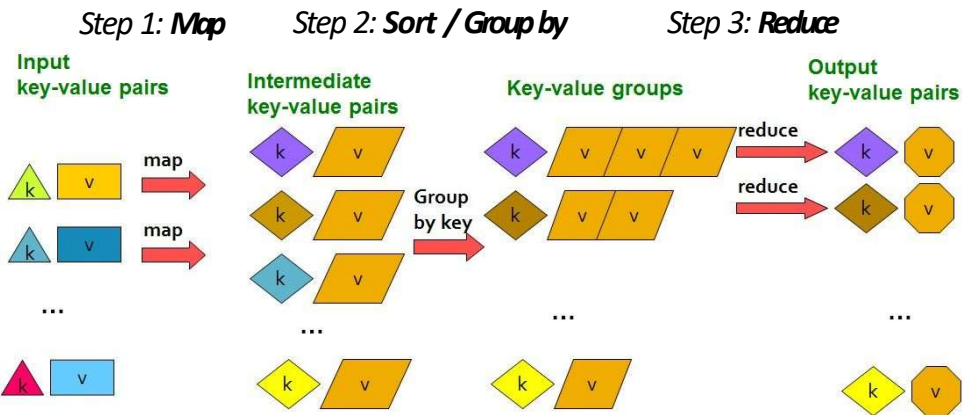


(Leskovec et al., 2014; <http://www.mmds.org/>)

42

What is MapReduce

Easy as 1, 2, 3!



(Leskovec et al., 2014; <http://www.mmds.org/>)

43

What is MapReduce

Map: $(k, v) \rightarrow (k', v')^*$
(Written by programmer)

Group by key: $(k_1', v_1'), (k_2', v_2'), \dots \rightarrow (k_1', (v_1', v', \dots)),$
(system handles) $(k_2', (v_1', v', \dots)), \dots$

Reduce: $(k', (v_1', v', \dots)) \rightarrow (k', v'')^*$
(Written by programmer)

44

Example: Word Count

```
tokenize(document) | sort | uniq -c
```

45

Example: Word Count

```
tokenize(document) | sort | uniq -c
```

Map: extract
what you
care about.

sort and
shuffle

Reduce:
aggregate,
summarize

46

Example: Word Count

The crew of the space shuttle Endeavor recently returned to Earth as ambassadors, harbingers of a new era of space exploration. Scientists at NASA are saying that the recent assembly of the Dextre bot is the first step in a long-term space-based man/machine partnership. "The work we're doing now -- the robotics we're doing - is what we're going to need

Big document

(Leskovec et al., 2014; <http://www.mmds.org/>)

47

Provided by the programmer

MAP:

Read input and produces a set of key-value pairs

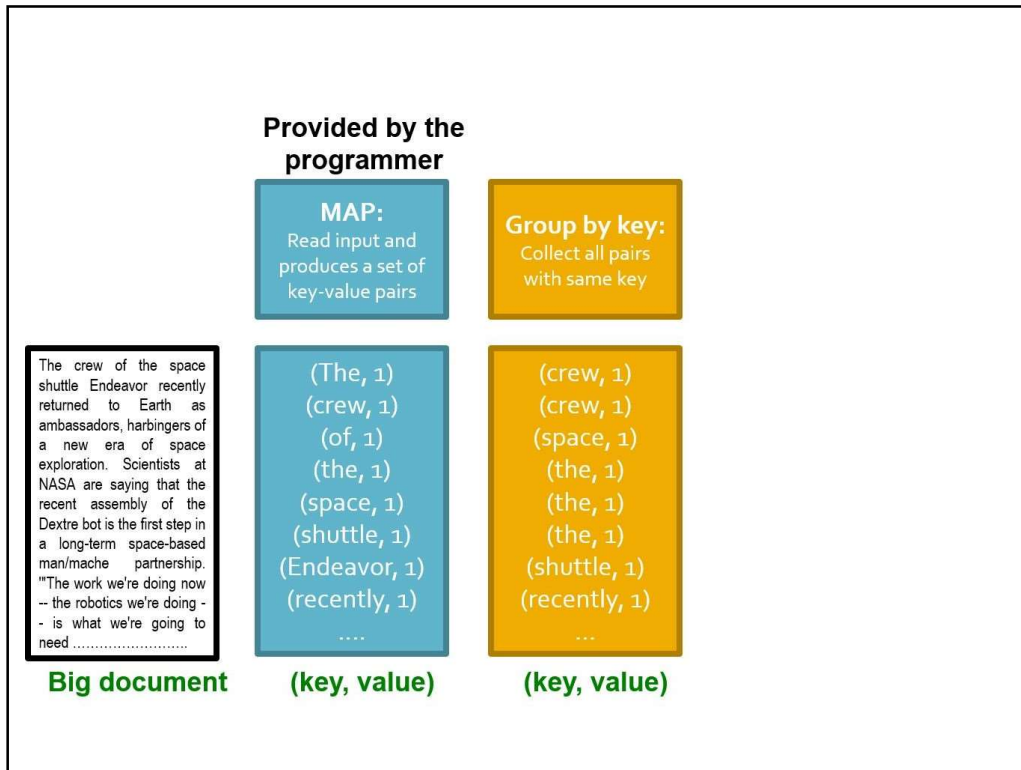
The crew of the space shuttle Endeavor recently returned to Earth as ambassadors, harbingers of a new era of space exploration. Scientists at NASA are saying that the recent assembly of the Dextre bot is the first step in a long-term space-based man/machine partnership. "The work we're doing now -- the robotics we're doing - is what we're going to need

Big document

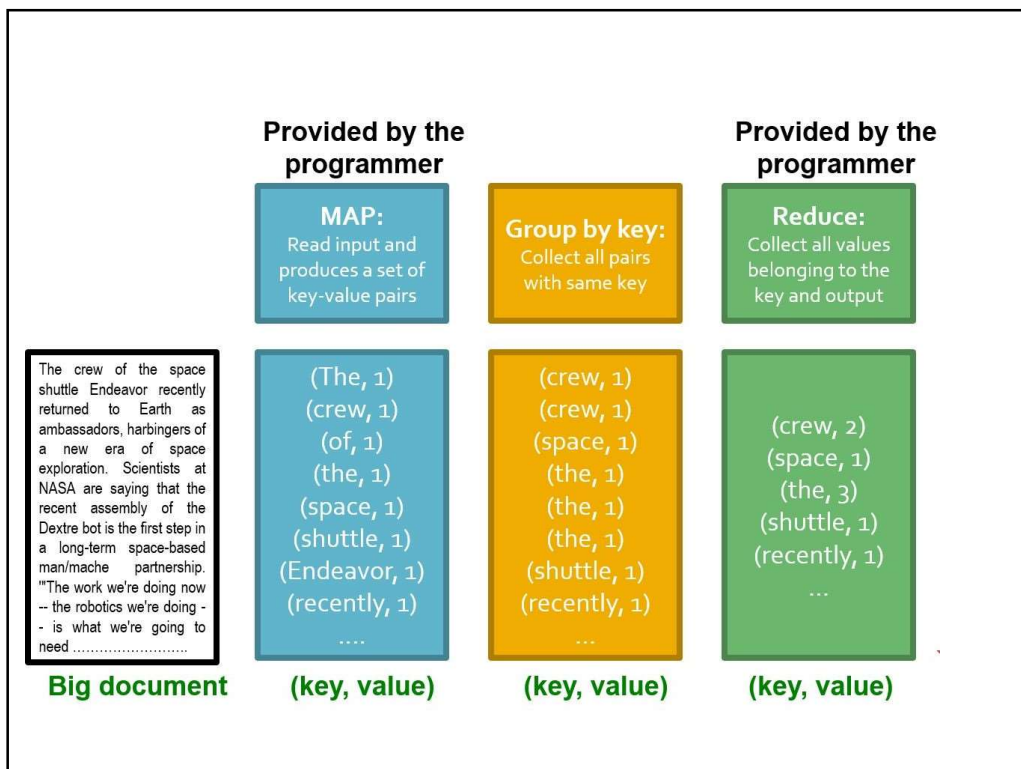
(The, 1)
(crew, 1)
(of, 1)
(the, 1)
(space, 1)
(shuttle, 1)
(Endeavor, 1)
(recently, 1)
.....

(key, value)

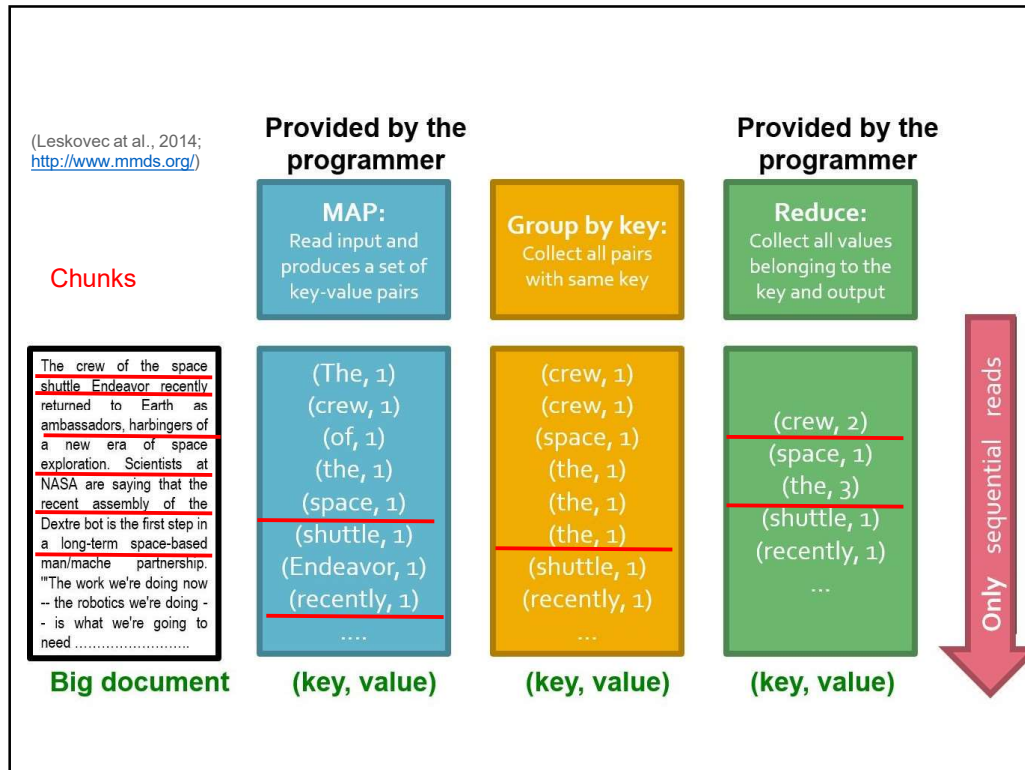
48



49



50



51

Example: Word Count

```
@abstractmethod
def map(k, v):
    pass

@abstractmethod
def reduce(k, vs):
    pass
```

52

Example: Word Count (v1)

```
def map(k, v):  
    for w in tokenize(v):  
        yield (w,1)
```

```
def reduce(k, vs):  
    return len(vs)
```

53

Example: Word Count (v1)

```
def map(k, v):  
    for w in tokenize(v):  
        yield (w,1)
```

```
def tokenize(s):  
    #simple version  
    return s.split(' ')
```

```
def reduce(k, vs):  
    return len(vs)
```

54

Example: Word Count (v2)

```
def map(k, v):  
    counts = dict()  
    for w in tokenize(v):
```

} counts each word within the chunk
(try/except is faster than
"if w in counts")

55

Example: Word Count (v2)

```
def map(k, v):  
    counts = dict()  
    for w in tokenize(v):  
        try:  
            counts[w] += 1  
        except KeyError:  
            counts[w] = 1  
    for item in counts.iteritems():  
        yield item
```

} counts each word within the chunk
(try/except is faster than
"if w in counts")

56

Example: Word Count (v2)

```
def map(k, v):
    counts = dict()
    for w in tokenize(v):
        try:
            counts[w] += 1
        except KeyError:
            counts[w] = 1
    for item in counts.iteritems():
        yield item
```

} counts each word within the chunk
(try/except is faster than
"if w in counts")



```
def reduce(k, vs):
    return (k, sum(vs) )
```

} sum of counts from different chunks

57

Distributed Architecture (Cluster)

Challenges for IO Cluster Computing

1. Nodes fail
1 in 1000 nodes fail a day
Duplicate Data (Distributed FS) 
2. Network is a bottleneck
Typically 1-10 Gb/s throughput
Bring computation to nodes, rather than data to nodes.
3. Traditional distributed programming is often ad-hoc and complicated
Stipulate a programming system that can easily be distributed

58

Distributed Architecture (Cluster)

Challenges for IO Cluster Computing

1. Nodes fail
1 in 1000 nodes fail a day
Duplicate Data (Distributed FS) ✓
2. Network is a bottleneck
Typically 1-10 Gb/s throughput
Bring computation to nodes, rather than data to nodes. (Sort and Shuffle) ✓
3. Traditional distributed programming is often ad-hoc and complicated
Stipulate a programming system that can easily be distributed

59

Distributed Architecture (Cluster)

Challenges for IO Cluster Computing

1. Nodes fail
1 in 1000 nodes fail a day
Duplicate Data (Distributed FS) ✓
2. Network is a bottleneck
Typically 1-10 Gb/s throughput
Bring computation to nodes, rather than data to nodes. (Sort and Shuffle) ✓
3. Traditional distributed programming is often ad-hoc and complicated **(Simply define a map and reduce)**
Stipulate a programming system that can easily be distributed ✓

60

Example: Relational Algebra

Select

Project

Union, Intersection, Difference

Natural Join

Grouping

61

Example: Relational Algebra

Select

Project

Union, Intersection, Difference

Natural Join

Grouping

62

Example: Relational Algebra

Select

$R(A_1, A_2, A_3, \dots)$, Relation R , Attributes A_*

return only those attribute tuples where condition C is true

63

Example: Relational Algebra

Select

$R(A_1, A_2, A_3, \dots)$, Relation R , Attributes A_*

return only those attribute tuples where condition C is true

```
def map(k, v): #v is list of attribute tuples: [(...), (...), ...]
    r = []
    for t in v:
        if t satisfies C:
            r += [(t, t)]
    return r
```

64

Example: Relational Algebra

Select

$R(A_1, A_2, A_3, \dots)$, Relation R , Attributes A_*

return only those attribute tuples where condition C is true

```
def map(k, v): #v is list of attribute tuples: [(...), (...), ...]
    r = []
    for t in v:
        if t satisfies C:
            r += [(t, t)]
    return r

def reduce(k, vs):
    r = []
    for each v in vs:
        r += [(k, v)]
    return r
```

65

Example: Relational Algebra

Select

$R(A_1, A_2, A_3, \dots)$, Relation R , Attributes A_*

return only those attribute tuples where condition C is true

```
def map(k, v): #v is list of attribute tuples
    for t in v:
        if t satisfies C:
            yield (t, t)

def reduce(k, vs):
    For each v in vs:
        yield (k, v)
```

66

Example: Relational Algebra

Natural Join

Given R_1 and R_2 return R_{join}

-- union of all pairs of tuples that match given attributes.

```
def map(k, v): #k \in {R1, R2}, v is (A, B) for R1, (B, C) for  
               R2 #B are matched attributes
```

67

Example: Relational Algebra

Natural Join

Given R_1 and R_2 return R_{join}

-- union of all pairs of tuples that match given attributes.

```
def map(k, v): #k \in {R1, R2}, v is (A, B) for R1, (B, C) for  
               R2 #B are matched attributes
```

```
    if k=='R1':  
        (a, b) = v  
        return (b, ('R1', a))  
    if k=='R2':  
        (b, c) = v  
        return (b, ('R2', c))
```

68

Example: Relational Algebra

Natural Join

Given R_1 and R_2 return R_{join}

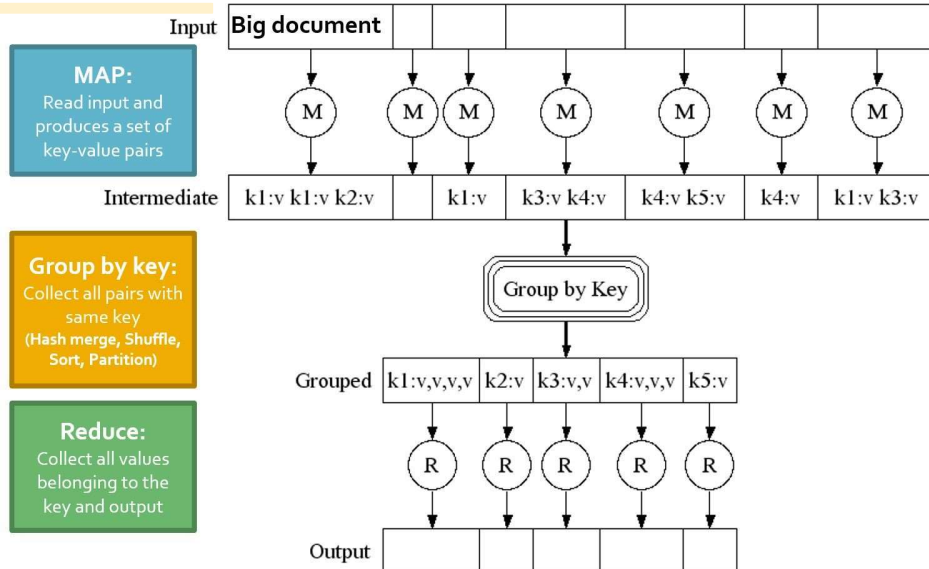
-- union of all pairs of tuples that match given attributes.

```
def map(k, v): #k \in {R1, R2}, v is (A, B) for R1, (B, C) for R2
    #B are matched attributes
    if k=='R1':
        (a, b) = v
        return (b, ('R1', a))
    if k=='R2':
        (b, c) = v
        return (b, ('R2', c))

def reduce(k, vs):
    r1, r2, rjn = [], [], []
    for (s, x) in vs: #separate rs
        if s == 'R1': r1.append(x)
        else: r2.append(x)
    for a in r1: #join as tuple
        for each c in r2:
            rjn += ('R_join', (a, k, c)) #k is b
    return rjn
```

69

Data Flow

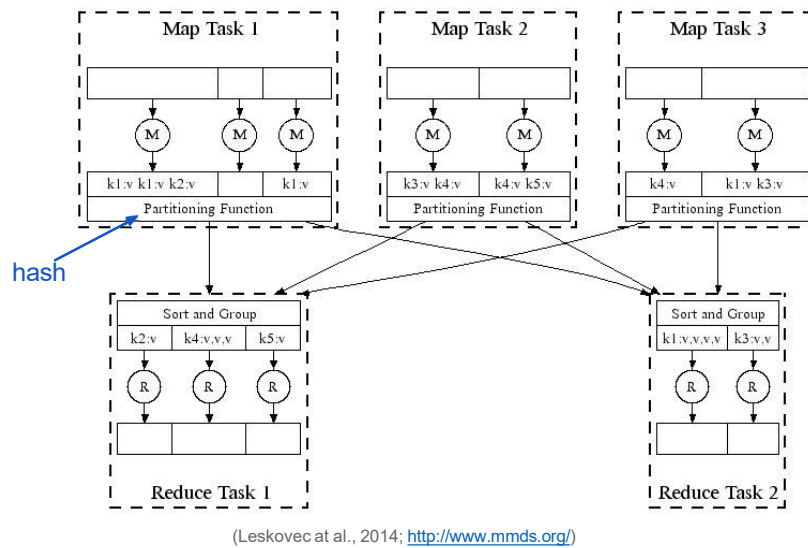


J. Leskovec, A. Rajaraman, J. Ullman: Mining of Massive Datasets, <http://www.mmds.org>

21

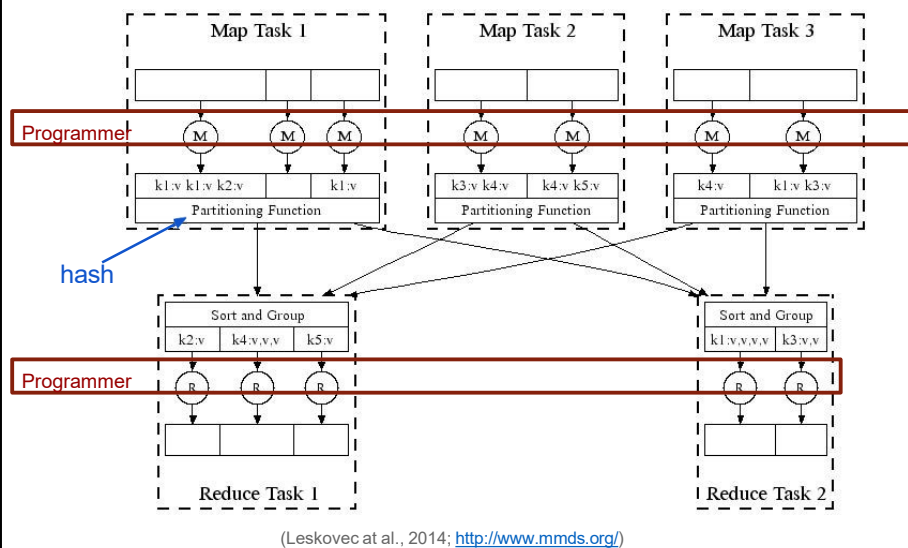
70

Data Flow



71

Data Flow



72

Data Flow

DFS \Rightarrow Map \Rightarrow Map's Local FS \Rightarrow Reduce \Rightarrow DFS

73

Data Flow

MapReduce system handles:

- Partitioning
- Scheduling map / reducer execution
- Group by key

- Restarts from node failures
- Inter-machine communication

74

Data Flow

DFS → MapReduce → DFS

- Schedule map tasks near physical storage of chunk
- Intermediate results stored locally
- Master / Name Node coordinates

75

Data Flow

DFS → MapReduce → DFS

- Schedule map tasks near physical storage of chunk
- Intermediate results stored locally
- Master / Name Node coordinates
 - Receives location of intermediate results and schedules with reducer
 - Checks nodes for failures and restarts when necessary
 - All map tasks on nodes must be completely restarted
 - Reduce tasks can pickup with reduce task failed

76

Data Flow

DFS → MapReduce → DFS

- Schedule map tasks near physical storage of chunk
- Intermediate results stored locally
- Master / Name Node coordinates
 - Task status: idle, in-progress, complete
 - Receives location of intermediate results and schedules with reducer
 - Checks nodes for failures and restarts when necessary
 - All map tasks on nodes must be completely restarted
 - Reduce tasks can pickup with reduce task failed

DFS → MapReduce → DFS → MapReduce → DFS

77

Data Flow

Skew: The degree to which certain tasks end up taking much longer than others.

Handled with:

- More reducers than reduce tasks
- More reduce tasks than nodes

78

Data Flow

Key Question: *How many Map and Reduce jobs?*

M: map tasks, *R*: reducer tasks

79

Data Flow

Key Question: *How many Map and Reduce jobs?*

M: map tasks, *R*: reducer tasks

Answer: 1) If possible, one chunk per map task, and

2) $M \gg |\text{nodes}| \approx |\text{cores}|$

(better handling of node failures, better load balancing)

3) $R \leq M$


(reduces number of parts stored in DFS)

80

Data Flow

□ Reduce Task

version 1: few reduce tasks
(same number of reduce tasks as nodes)




Reduce tasks represented by
time to complete task
(some tasks take much longer)

81

Data Flow


□ Reduce Task

version 1: few reduce tasks
(same number of reduce tasks as nodes)



Reduce tasks represented by
time to complete task
(some tasks take much longer)

version 2: more reduce tasks
(more reduce tasks than nodes)




Reduce tasks represented by
time to complete task
(some tasks take much longer)

82

Data Flow


version 1: few reduce tasks
(same number of reduce tasks as nodes)



Reduce tasks represented by
time to complete task
(some tasks take much longer)

□ Reduce Task


version 2: more reduce tasks
(more reduce tasks than nodes)



Reduce tasks represented by
time to complete task
(some tasks take much longer)

Last task
completed

Can
redistribute
these tasks to
other nodes



(the last task now completes
much earlier)

83

Communication Cost Model

How to assess performance?

- (1) Computation: Map + Reduce + System Tasks
- (2) Communication: Moving (key, value) pairs

84

Communication Cost Model

How to assess performance?

- (1) Computation: Map + Reduce + System Tasks
- (2) Communication: Moving (key, value) pairs

Ultimate Goal: wall-clock Time.



85

Communication Cost Model

How to assess performance?

- (1) **Computation: Map + Reduce + System Tasks**

- Mappers and reducers often single pass $O(n)$ within node
- System: sort the keys is usually most expensive
- Even if map executes on same node, disk read usually dominates
- In any case, can add more nodes

Ultimate Goal: wall-clock Time.



86

Communication Cost Model

How to assess performance?

(1) Computation: Map + Reduce + System Tasks

(2) Communication: Moving key, value pairs

Often dominates computation.

- Connection speeds: 1-10 gigabits per sec;
- HD read: 50-150 gigabytes per sec
- Even reading from disk to memory typically takes longer than operating on the data.

87

Communication Cost Model

How to assess performance?

Communication Cost = input size +
(sum of size of all map-to-reducer files)

(2) Communication: Moving key, value pairs

Often dominates computation.

- Connection speeds: 1-10 gigabits per sec;
- HD read: 50-150 gigabytes per sec
- Even reading from disk to memory typically takes longer than operating on the data.

88

Communication Cost Model

How to assess performance?

Communication Cost = input size +
(sum of size of all map-to-reducer files)

(2) Communication: Moving key, value pairs

Often dominates computation.

- Connection speeds: 1-10 gigabits per sec;
H/D read: 50-150 gigabytes per sec
- Even reading from disk to memory typically takes longer than operating on the data.
- Output from reducer ignored because it's either small (finished summarizing data) or being passed to another mapreduce job.

89

Communication Cost: Natural Join

R, S: Relations (Tables) $R(A, B) \bowtie S(B, C)$

Communication Cost = input size +
(sum of size of all map-to-reducer files)

DFS → Map → LocalFS → Network → Reduce → DFS → ?

90

Communication Cost: Natural Join

R, S: Relations (Tables) $R(A, B) \bowtie S(B, C)$

Communication Cost = input size +
(sum of size of all map-to-reducer files)

```
def map(k, v):
    if k=="R1":
        (a, b) = v
        yield (b, (R1, a))
    if k=="R2":
        (b, c) = v
        yield (b, (R2, c))

def reduce(k, vs):
    r1, r2 = [], []
    for (rel, x) in vs: #separate rs
        if rel == 'R': r1.append(x)
        else: r2.append(x)
    for a in r1: #join as tuple
        for each c in r2:
            yield (Rjoin, (a, k, c)) #k is
```

91

Communication Cost: Natural Join

R, S: Relations (Tables) $R(A, B) \bowtie S(B, C)$

Communication Cost = input size +
(sum of size of all map-to-reducer files)

= |R₁| + |R₂| + (|R₁| + |R₂|)

= O(|R₁| + |R₂|)

```
def map(k, v):
    if k=="R1":
        (a, b) = v
        yield (b, (R1, a))
    if k=="R2":
        (b, c) = v
        yield (b, (R2, c))

def reduce(k, vs):
    r1, r2 = [], []
    for (rel, x) in vs: #separate rs
        if rel == 'R': r1.append(x)
        else: r2.append(x)
    for a in r1: #join as tuple
        for each c in r2:
            yield (Rjoin, (a, k, c)) #k is
```

92

MapReduce: Final Considerations

- Performance Refinements:
 - Combiners (like word count version 2 but done via reduce)
 - Run reduce right after map from same node before passing to reduce (MapTask can execute)
 - Reduces communication cost
- Requires commutative and associative reducer function.

93

MapReduce: Final Considerations

- Performance Refinements:
 - Combiners (like word count version 2 but done via reduce)
 - Run reduce right after map from same node before passing to reduce (MapTask can execute)
 - Reduces communication cost
 - Backup tasks (aka speculative tasks)
 - Schedule multiple copies of tasks when close to the end to mitigate certain nodes running slow.
 - Override partition hash function to organize data
E.g. instead of `hash(url)` use `hash(hostname(url))`

94